

Abhay Singh

contact

✉ as2626@cornell.edu
☎ (404) 353-0477

in [linkedin.com/in/as2626](https://www.linkedin.com/in/as2626)
github.com/as2626

education

Cornell University, Ithaca, NY

M.S. in Computer Science

Aug 2021 – May 2023

B.S. in Computer Science with Honors, *Summa Cum Laude*



Aug 2018 – May 2021



coursework

† = teaching
assistant

CS 6241: Data Science Numerics	ECE 6970: Statistical Distances	CS 6670: Computer Vision
ORIE 6510: Probability	CS 4780: Machine Learning †	CS 4820: Algorithms †
MATH 4130: Analysis I (Honors)	MATH 4315: Linear Algebra	CS 4850: Math Foundations

publications

Combining Label Propagation and Simple Models Out-performs GNNs   (ICLR 2021)
Qian Huang, Horace He, Abhay Singh, Ser-Nam Lim, and Austin Benson

Better Set Representations For Relational Reasoning   (NeurIPS 2020)
Qian Huang, Horace He, Abhay Singh, Yan Zhang, Ser-Nam Lim, and Austin Benson

experience

Yext, New York, NY

Software Engineering Intern

May 2020 – Aug 2020

- Designed and integrated static code analysis tool used firm-wide on over 80% of codebase to scan vulnerable Java code at compile-time
- Wrote multi-threaded Golang script to determine unprotected customer apps that downloads and parses terabytes of API log data on-the-fly via AWS S3, and makes remote-procedure calls to fetch app data by API key; improved performance by 4x relative to previous solution
- Integrated webhooks to automate modification of company repository permissions using Github's REST API, notifying teams automatically via Slack and email

Morgan Stanley, New York, NY

Technology Summer Analyst

June 2019 – Aug 2019

- Architected and implemented end-to-end data pipeline to process and analyze over 800,000,000 entries of financial data daily with highly optimized, parallelized Python scripts, using NumPy and Pandas
- Reduced mainframe consumption by 90%, from 5000 to 500 CPU-seconds, saving tens of millions of dollars in annual costs
- Created and deployed firm-wide DevOps web tool to analyze large text-based datasets


Cornell Unmanned Air Systems, Ithaca, NY

Computer Vision and DevOps Lead

Oct 2018 – May 2020

- Designed and implemented custom object detection and classification model (Mask R-CNN variant with multi-head output) in multi-task learning setting on collected aerial imagery dataset, in PyTorch
- Lead all computer vision tasks on team, with individual efforts directly increasing classification task accuracy by 32% and object detection mAP IoU by over 80%


projects

1-Lipschitz Deep Equilibrium Models 

- Presented method to enforce the uniqueness and existence of a fixed point solution in a neural network representing an iterative solver, which is done by constraining 1-Lipschitzness of the model

Few-Shot Clustering Instance Segmentation (FS-CIS) Net 

- Designed novel model architecture to perform proposal-free few-shot instance segmentation in autonomous driving scenarios, speeding up inference with comparable performance to existing methods

Continual Learning with Lottery Tickets 

- Proposed and demonstrated effectiveness of novel training scheme to resist catastrophic forgetting, a phenomena in which a model overfits to the most recently seen data in a multi-task learning setting

Xi Compiler

- Wrote optimized compiler in Scala for language Xi, in team of 4; approximately 10,000 lines of code
- Includes lexing, parsing, type-checking, intermediate code generation, various optimizations including dataflow analysis, and emitting assembly instructions with non-trivial register allocation

languages & technologies

Python, Java, OCaml, Scala, Go, Julia, C/C++, Bash, JavaScript, SQL
PyTorch, Keras/TensorFlow, Git, Jupyter, Docker, Bazel, Gradle, Terraform